

Contents

Cloud Architecture Design Document	1
Executive Summary	1
System Overview	1
Cloud Deployment Architecture	2
Microservices Architecture	3
Service Interconnections & Communication Mechanisms	11
Design Rationale & Decision Log	15
Infrastructure as Code	19
Scalability & Resilience	20
Observability	20
Security Considerations	20
Cost Optimization	21
Future Enhancements	21
Conclusion	21

Cloud Architecture Design Document

Executive Summary

This document provides a comprehensive overview of a production-grade, cloud-native e-commerce platform built using microservices architecture with multi-cloud deployment across AWS and GCP. The system demonstrates modern DevOps practices including Infrastructure as Code, GitOps, event-driven architecture, and comprehensive observability.

System Overview

Business Context

The platform is a scalable e-commerce solution that enables:

- **Customer Operations:** User registration, authentication, and profile management
- **Product Management:** Dynamic product catalog with search and filtering
- **Order Processing:** Real-time order creation, payment processing, and order tracking
- **Notifications:** Multi-channel customer notifications (Email, SMS, Push)
- **Analytics:** Real-time business intelligence and reporting

Technical Goals

1. **High Availability:** 99.9% uptime with multi-AZ deployment
2. **Scalability:** Auto-scaling from 2 to 100+ concurrent users
3. **Event-Driven Architecture:** Asynchronous processing with Kafka
4. **Real-Time Analytics:** Stream processing with Apache Flink
5. **Infrastructure as Code:** Reproducible environments with Terraform
6. **GitOps Deployment:** Automated deployments with ArgoCD
7. **Observability:** Comprehensive monitoring with Prometheus/Grafana
8. **Security:** Zero-trust architecture with IAM and network policies

Technology Stack

- **Container Orchestration:** Kubernetes (EKS on AWS, GKE on GCP)
- **Microservices Framework:** Python Flask, Spring Boot (Analytics)
- **Message Broker:** Apache Kafka (AWS MSK)
- **Databases:** PostgreSQL (RDS), DynamoDB, Redis
- **Stream Processing:** Apache Flink on GCP Dataflow
- **Infrastructure:** Terraform
- **CI/CD:** ArgoCD, GitHub Actions

- **Monitoring:** Prometheus, Grafana, EFK Stack
-

Cloud Deployment Architecture

High-Level Architecture Diagram

INTERNET / USERS □ (HTTPS)

AWS CLOUD (Primary)

VPC (10.0.0.0/16) - Multi-AZ Architecture

1. **Public Subnets** (us-east-1a, us-east-1b)
 - AWS Application Load Balancer (Layer 7)
 - Ingress Controller (NGINX)
2. **Amazon EKS Cluster** (Kubernetes 1.27)
 - **Microservices Pods:**
 - API Gateway (2-10 replicas with HPA)
 - User Service (2 replicas static)
 - Product Service (2 replicas static)
 - Order Service (2-8 replicas with HPA)
 - Payment Service (2 replicas static)
 - Notification Service (1 replica)
 - **Management:**
 - Horizontal Pod Autoscaler (HPA)
 - Prometheus (Metrics Collection)
 - ArgoCD (GitOps Deployment)
3. **Private Subnets** (us-east-1a, us-east-1b)
 - RDS PostgreSQL (Multi-AZ) - Users, Orders, Payments
 - DynamoDB (On-demand) - Product Catalog
 - Redis (ElastiCache) - Caching Layer
 - AWS MSK (Kafka) - Event Streaming
 - orders topic (3 partitions, replication factor 2)
 - analytics-results topic (3 partitions, replication factor 2)
4. **Additional AWS Services**
 - S3 (Object Storage)
 - Lambda (Serverless Functions)
 - ECR (Container Registry)
 - Route 53 (DNS)
 - CloudWatch (Monitoring)

□ (Cross-Cloud Kafka Connection)

GOOGLE CLOUD (Analytics)

GCP Dataproc Cluster - Apache Flink Analytics Job - Kafka Consumer (orders topic) - Stream Processing (Window Aggregation) - Kafka Producer (analytics-results topic) - Google Cloud Storage (GCS) - Flink checkpoints - Analytics results - Data lake storage

Multi-Cloud Strategy

Primary Cloud: AWS - EKS (Elastic Kubernetes Service): Primary container orchestration platform - **RDS PostgreSQL:** Relational data for users, orders, payments - **DynamoDB:** NoSQL storage for product catalog - **MSK (Managed Streaming for Kafka):** Event streaming backbone - **S3:** Object storage for static assets and backups - **Lambda:** Serverless functions for event processing - **ALB:** Application Load Balancer for traffic distribution - **ECR:** Container image registry - **Route 53:** DNS management - **CloudWatch:** Basic monitoring and logging

Secondary Cloud: GCP - Dataproc: Managed Spark/Flink cluster for stream processing - **Cloud Storage (GCS):** Data lake and checkpoints - **Cloud Pub/Sub:** Alternative messaging (future) - **BigQuery:** Data warehousing (future)

Deployment Rationale

Why AWS as Primary Cloud?

1. **Mature Kubernetes Platform:** EKS is battle-tested with excellent AWS service integration
2. **Rich Database Options:** RDS and DynamoDB cover all data storage needs
3. **Managed Kafka:** MSK eliminates operational overhead of running Kafka clusters
4. **Cost Effectiveness:** Competitive pricing with reserved instances and savings plans
5. **Enterprise Support:** Extensive documentation and community support
6. **Compliance:** SOC 2, HIPAA, PCI-DSS certifications

Why GCP for Analytics?

1. **Best-in-Class Analytics:** Dataproc with Flink provides superior stream processing
2. **Workload Isolation:** Separates compute-intensive analytics from transactional workloads
3. **Cost Optimization:** Pay only for analytics compute when needed
4. **Data Processing Excellence:** GCP's heritage with Hadoop ecosystem
5. **Future Integration:** Easy path to BigQuery and ML services

Multi-Cloud Benefits

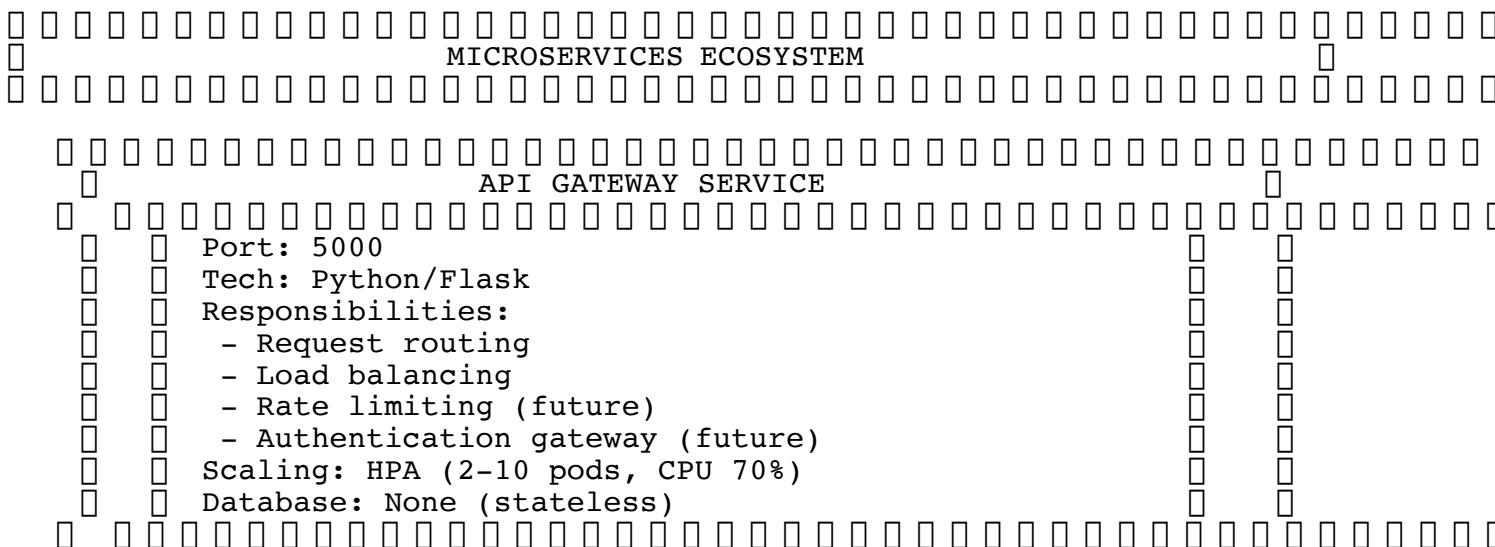
- **Vendor Lock-in Avoidance:** Not dependent on single cloud provider
 - **Best-of-Breed Services:** Use optimal service from each provider
 - **Risk Mitigation:** Disaster recovery across cloud providers
 - **Skills Development:** Team gains multi-cloud expertise
 - **Negotiation Power:** Better pricing negotiations with multiple options
-

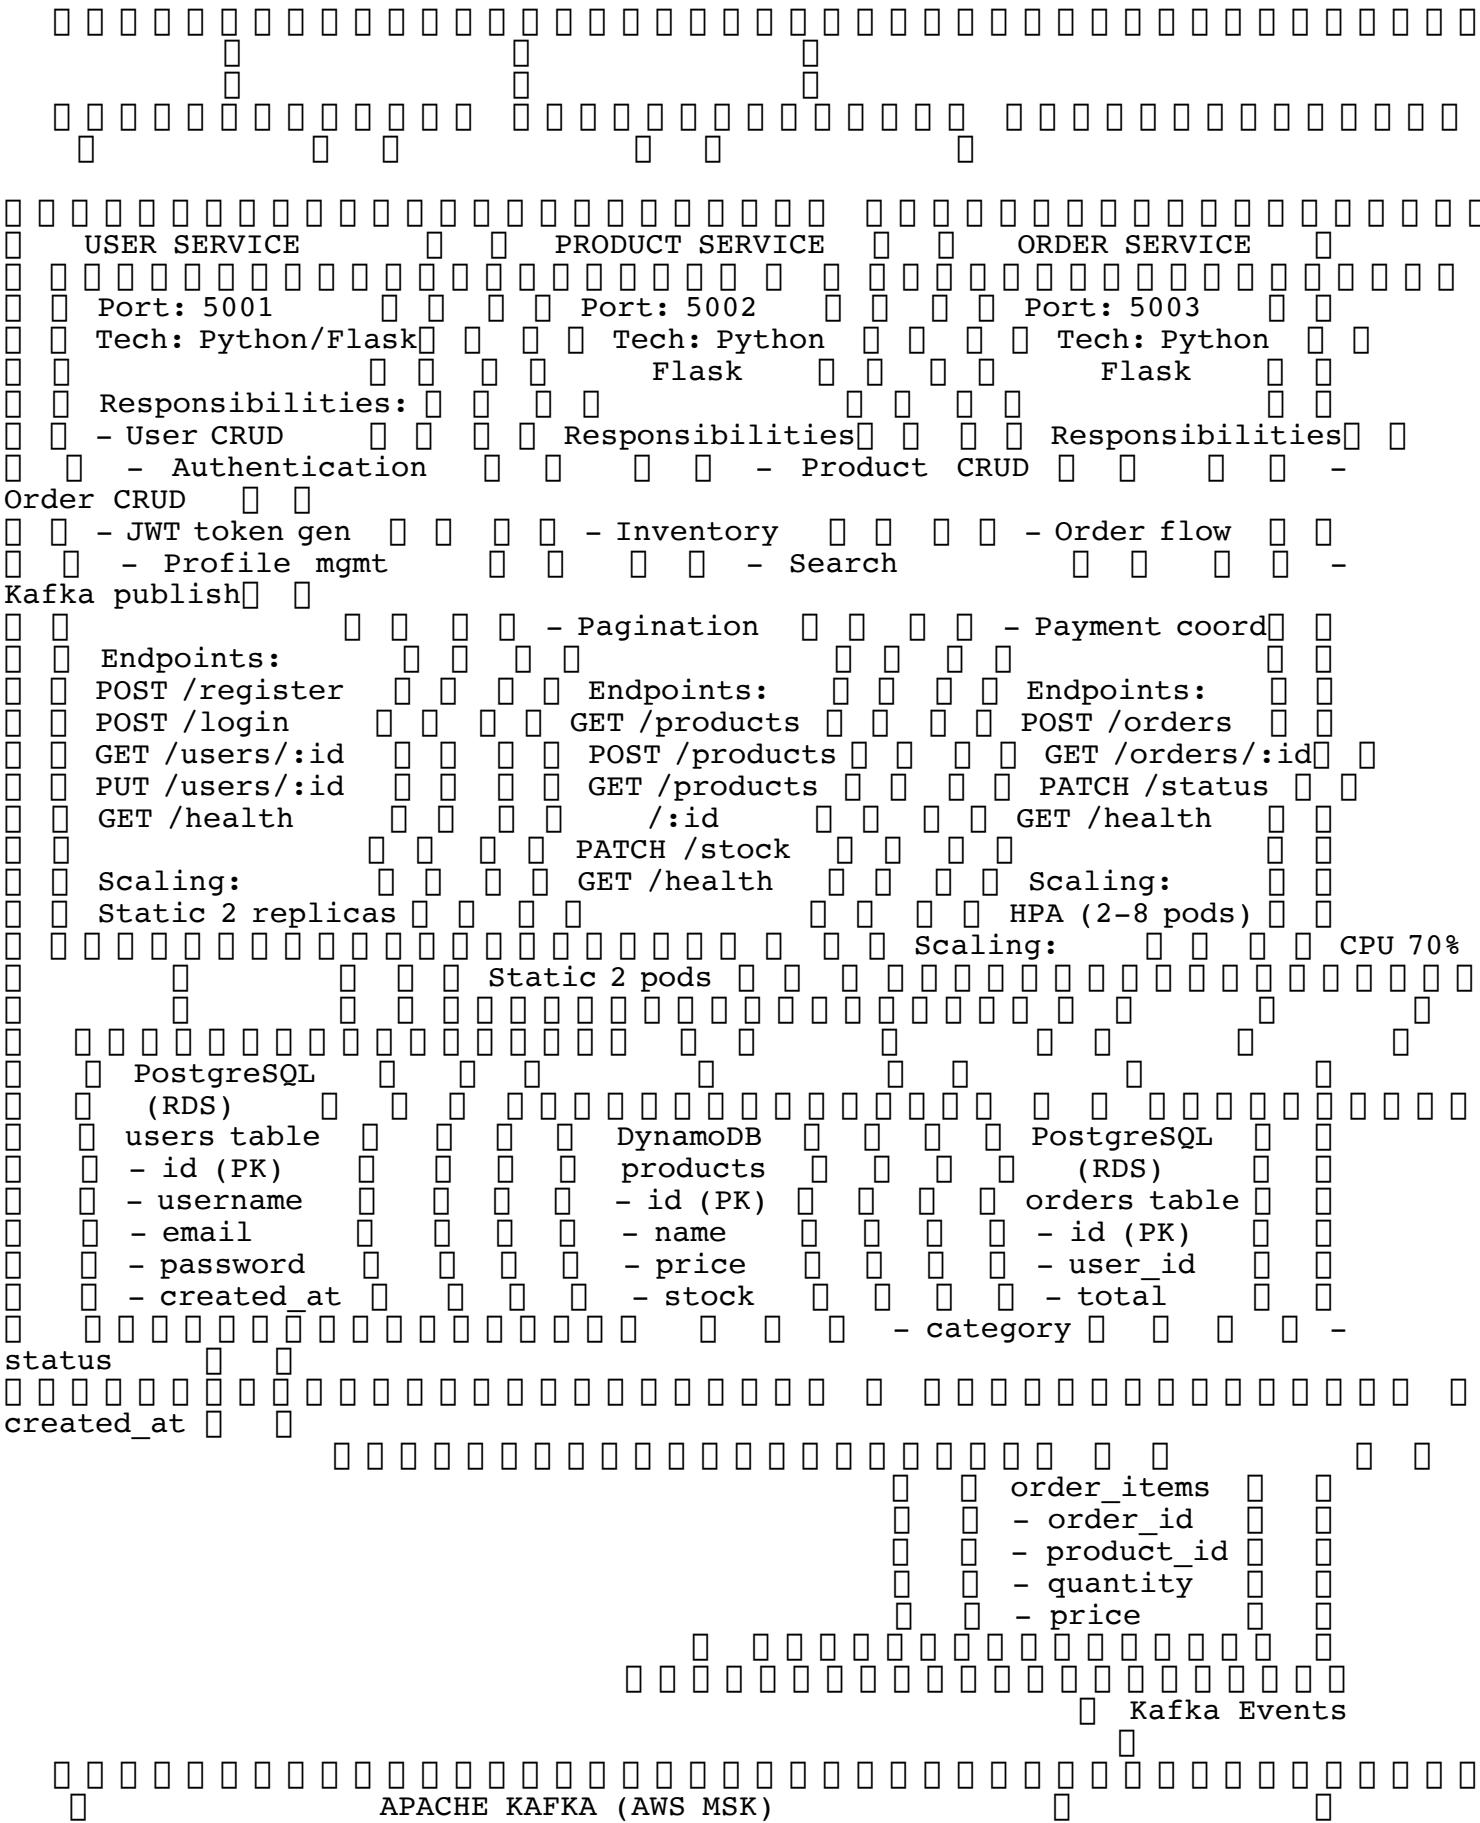
Microservices Architecture

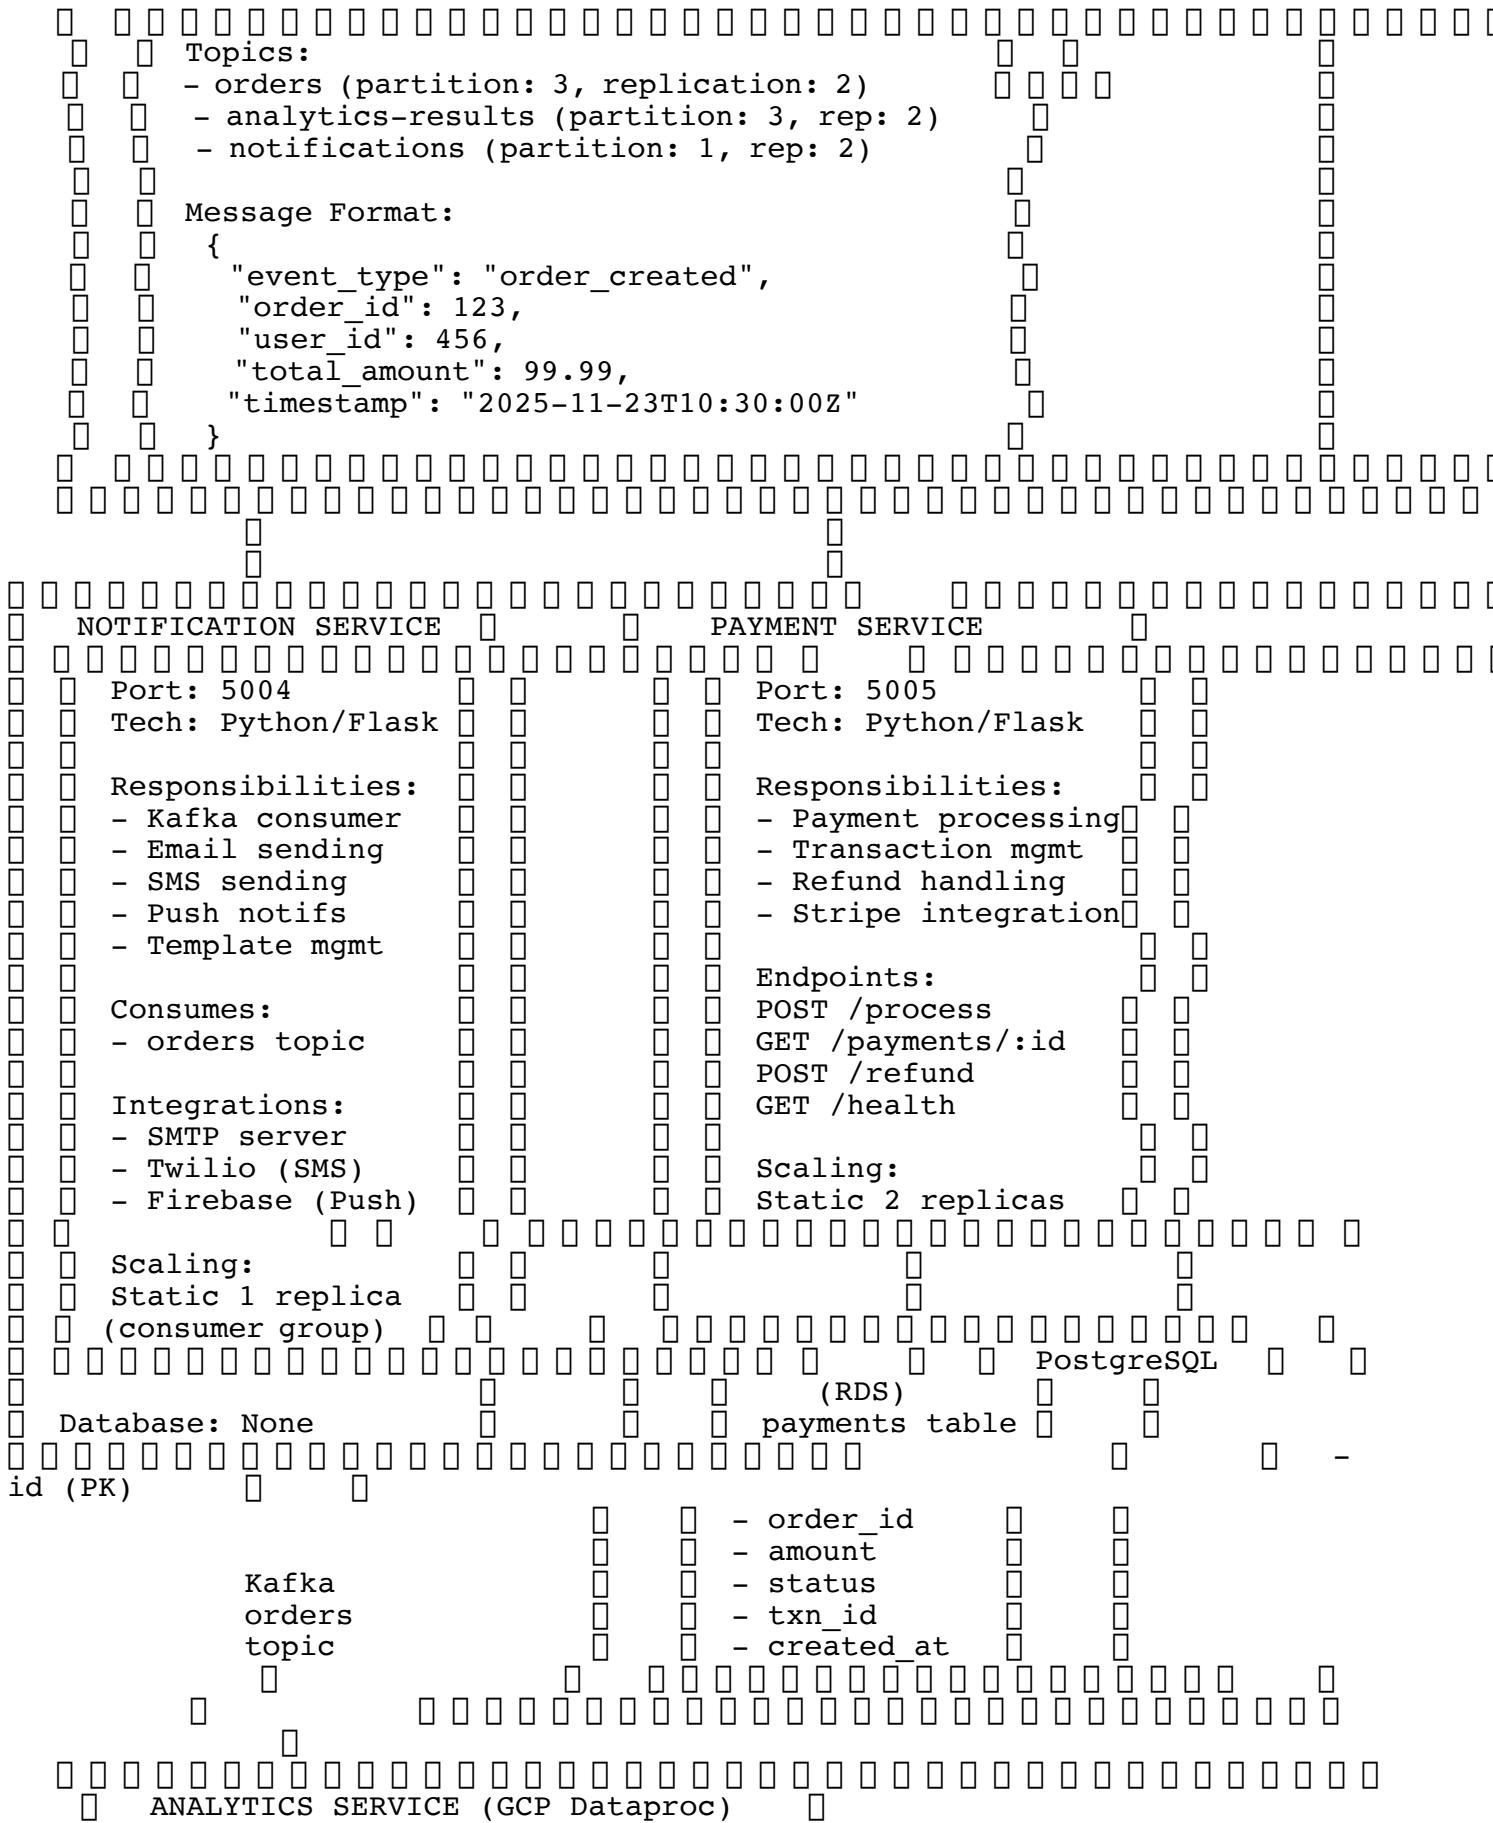
Architecture Principles

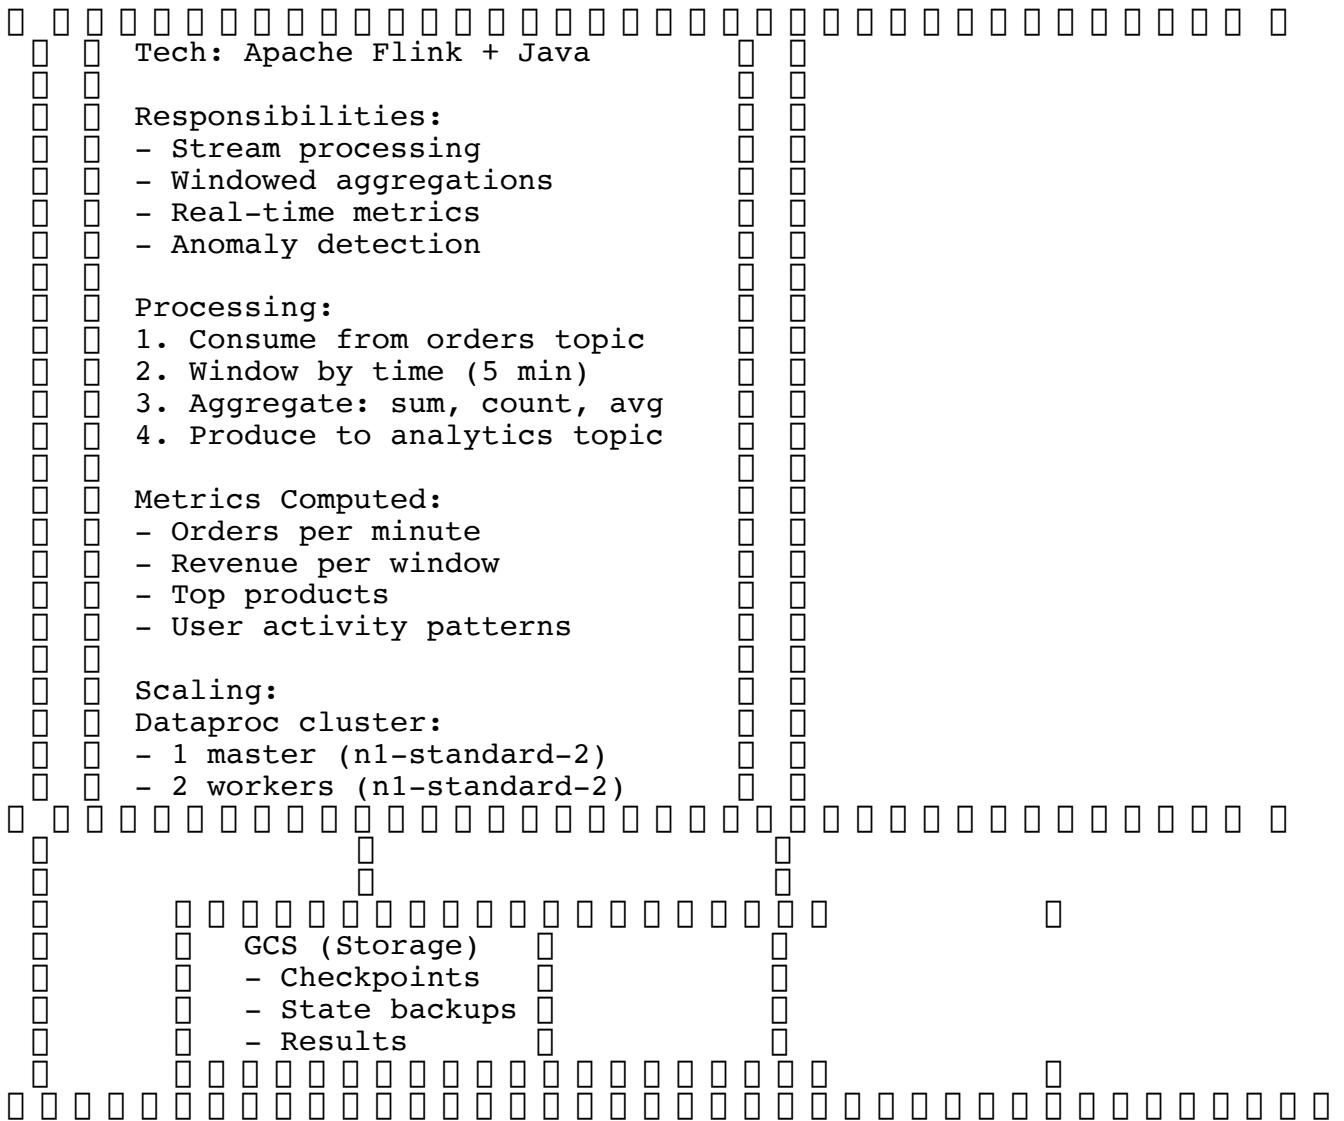
1. **Single Responsibility:** Each service owns a specific business capability
2. **Loose Coupling:** Services communicate via well-defined APIs
3. **High Cohesion:** Related functionality grouped within services
4. **Independent Deployment:** Services can be deployed independently
5. **Technology Diversity:** Different tech stacks per service needs
6. **Data Ownership:** Each service owns its database

Detailed Microservices Diagram









Service Inventory & Detailed Responsibilities

1. API Gateway Service

- **Port:** 5000
- **Technology:** Python 3.11, Flask 3.0, Gunicorn
- **Storage:** None (stateless)
- **Primary Functions:**
 - Route HTTP requests to appropriate microservices
 - Load distribution across service instances
 - Request/response transformation
 - Error handling and standardization
 - Health check aggregation
- **Communication Pattern:**
 - Inbound: HTTP REST from Internet via AWS ALB
 - Outbound: HTTP REST to all internal services
- **Scaling Strategy:**
 - HPA based on CPU (70%) and memory (80%)
 - Min: 2 replicas, Max: 10 replicas
 - Scale up: 100% increase every 15s or 2 pods every 15s (whichever is higher)
 - Scale down: 50% decrease every 60s with 300s stabilization window

- **Resource Allocation:**
 - Requests: 128Mi memory, 100m CPU
 - Limits: 256Mi memory, 200m CPU

2. User Service

- **Port:** 5001
- **Technology:** Python 3.11, Flask 3.0, SQLAlchemy 2.0, JWT
- **Storage:** PostgreSQL 15 (RDS Multi-AZ)
- **Primary Functions:**
 - User registration with password hashing (bcrypt)
 - Authentication and JWT token generation
 - User profile CRUD operations
 - Password reset functionality
 - Session management

- **Database Schema:**

```
CREATE TABLE users (
    id SERIAL PRIMARY KEY,
    username VARCHAR(80) UNIQUE NOT NULL,
    email VARCHAR(120) UNIQUE NOT NULL,
    password_hash VARCHAR(255) NOT NULL,
    full_name VARCHAR(120),
    phone VARCHAR(20),
    address TEXT,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

- **API Endpoints:**
 - POST /api/users/register - Create new user
 - POST /api/users/login - Authenticate and get JWT
 - GET /api/users/:id - Get user profile
 - PUT /api/users/:id - Update user profile
 - DELETE /api/users/:id - Soft delete user
 - GET /health - Health check

- **Scaling Strategy:** Static 2 replicas (low variance in load)

- **Resource Allocation:**

- Requests: 128Mi memory, 100m CPU
- Limits: 256Mi memory, 200m CPU

3. Product Service

- **Port:** 5002
- **Technology:** Python 3.11, Flask 3.0, Boto3 (DynamoDB SDK)
- **Storage:** DynamoDB (On-demand billing)
- **Primary Functions:**
 - Product catalog management
 - Inventory tracking
 - Product search and filtering
 - Category management
 - Stock level monitoring

- **Database Schema** (DynamoDB):

Table: products
 Partition Key: id (Number)
 Attributes:

- name (String)
- description (String)
- price (Number)
- category (String)
- stock_quantity (Number)
- image_url (String)
- sku (String)
- created_at (String - ISO 8601)
- updated_at (String - ISO 8601)

Global Secondary Indexes:

- category-index (category as partition key)
- sku-index (sku as partition key)

- **API Endpoints**:

- GET /api/products - List products (paginated)
- GET /api/products/:id - Get product details
- POST /api/products - Create product (admin)
- PUT /api/products/:id - Update product (admin)
- DELETE /api/products/:id - Delete product (admin)
- PATCH /api/products/:id/stock - Update stock
- GET /health - Health check

- **Scaling Strategy**: Static 2 replicas

- **Resource Allocation**:

- Requests: 256Mi memory, 250m CPU
- Limits: 512Mi memory, 500m CPU

4. Order Service

- **Port**: 5003
- **Technology**: Python 3.11, Flask 3.0, SQLAlchemy 2.0, Kafka-Python
- **Storage**: PostgreSQL 15 (RDS Multi-AZ)
- **Primary Functions**:

- Order creation and validation
- Order status management workflow
- Integration with payment service
- Integration with product service (stock check)
- Publish order events to Kafka
- Order history tracking

- **Database Schema**:

```
CREATE TABLE orders (
  id SERIAL PRIMARY KEY,
  user_id INTEGER NOT NULL,
  total_amount DECIMAL(10, 2) NOT NULL,
  status VARCHAR(50) DEFAULT 'pending',
  shipping_address TEXT,
  payment_id VARCHAR(100),
  created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
```

```

    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE order_items (
    id SERIAL PRIMARY KEY,
    order_id INTEGER REFERENCES orders(id),
    product_id INTEGER NOT NULL,
    product_name VARCHAR(200) NOT NULL,
    quantity INTEGER NOT NULL,
    price DECIMAL(10, 2) NOT NULL
);

```

- **Order Status Flow:**

1. pending □ Order created, awaiting payment
2. processing □ Payment confirmed
3. shipped □ Order dispatched
4. delivered □ Order completed
5. cancelled □ Order cancelled

- **API Endpoints:**

- POST /api/orders - Create order
- GET /api/orders/:id - Get order details
- GET /api/orders/user/:userId - Get user orders
- PATCH /api/orders/:id/status - Update status
- DELETE /api/orders/:id - Cancel order
- GET /health - Health check

- **Kafka Integration:**

- Produces to `orders` topic on order creation
- Event payload includes order details for downstream processing

- **Scaling Strategy:**

- HPA based on CPU (70%) and memory (80%)
- Min: 2 replicas, Max: 8 replicas
- Scale up: 100% increase every 15s or 2 pods every 15s
- Scale down: 50% decrease every 60s

- **Resource Allocation:**

- Requests: 128Mi memory, 100m CPU
- Limits: 256Mi memory, 200m CPU

5. Payment Service

- **Port:** 5005
- **Technology:** Python 3.11, Flask 3.0, SQLAlchemy 2.0, Stripe SDK
- **Storage:** PostgreSQL 15 (RDS Multi-AZ)
- **Primary Functions:**

- Payment processing integration
- Transaction management
- Refund handling
- Payment method management
- PCI compliance (via Stripe)

- **Database Schema:**

```

CREATE TABLE payments (
    id SERIAL PRIMARY KEY,
    order_id INTEGER UNIQUE NOT NULL,
    amount DECIMAL(10, 2) NOT NULL,
    payment_method VARCHAR(50),
    transaction_id VARCHAR(100) UNIQUE,
    status VARCHAR(50) DEFAULT 'pending',
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

```

- **API Endpoints:**

- POST /api/payments/process - Process payment
- GET /api/payments/:id - Get payment details
- POST /api/payments/refund - Issue refund
- GET /api/payments/order/:orderId - Get payment by order
- GET /health - Health check

- **Scaling Strategy:** Static 2 replicas

- **Resource Allocation:**

- Requests: 128Mi memory, 100m CPU
- Limits: 256Mi memory, 200m CPU

6. Notification Service

- **Port:** 5004

- **Technology:** Python 3.11, Flask 3.0, Kafka-Python, SMTP, Twilio

- **Storage:** None (stateless)

- **Primary Functions:**

- Consume order events from Kafka
- Send email notifications (SMTP/SendGrid)
- Send SMS notifications (Twilio)
- Send push notifications (Firebase)
- Template management
- Notification retry logic

- **Notification Types:**

- Order confirmation
- Order status updates
- Payment confirmation
- Shipping notifications
- Promotional emails (future)

- **API Endpoints:**

- POST /api/notifications/send - Manual notification
- GET /health - Health check

- **Kafka Integration:**

- Consumes from orders topic
- Consumer group: notification-service-group
- Auto-commit: false (manual offset management)

- **Scaling Strategy:** Static 1 replica (single consumer per partition)

- **Resource Allocation:**

- Requests: 256Mi memory, 250m CPU

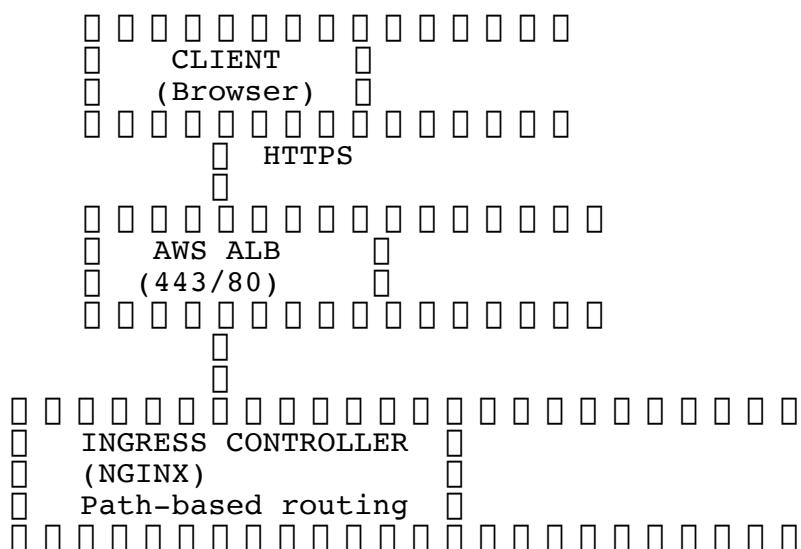
- Limits: 512Mi memory, 500m CPU

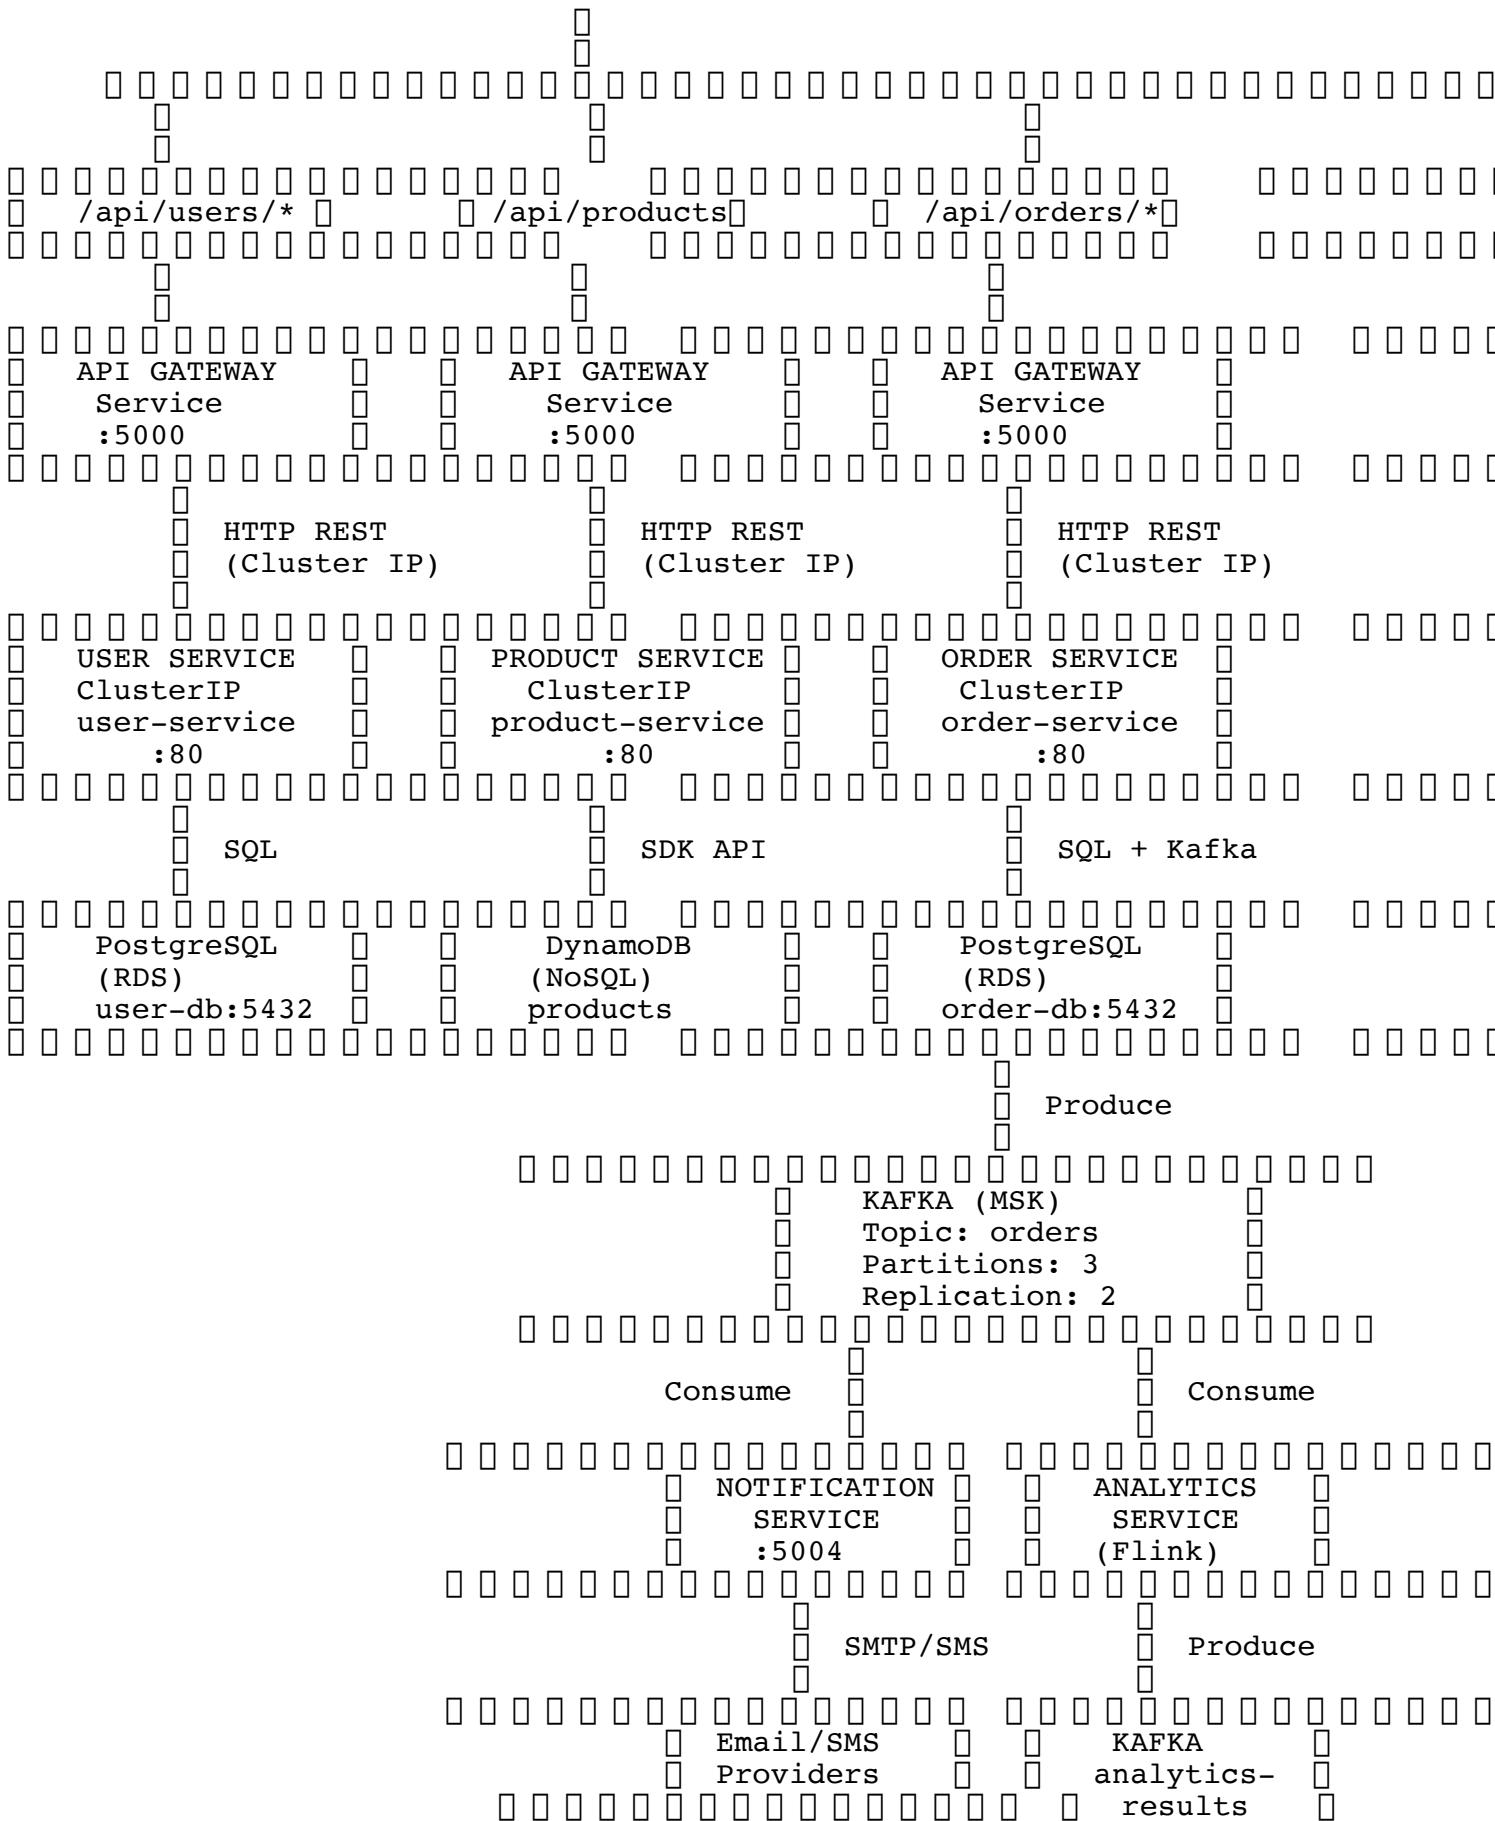
7. Analytics Service (Flink)

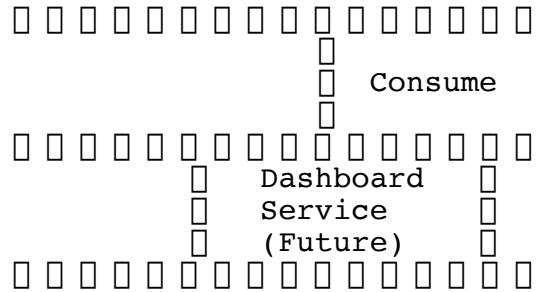
- **Technology:** Apache Flink 1.17, Java 11, Kafka Connector
 - **Infrastructure:** GCP Dataproc cluster
 - **Primary Functions:**
 - Real-time stream processing
 - Windowed aggregations (tumbling, sliding)
 - Event time processing
 - Stateful computations
 - Anomaly detection
 - **Processing Pipeline:**
 1. **Source:** Kafka consumer (orders topic)
 2. **Transform:** Parse JSON, extract fields
 3. **Window:** 5-minute tumbling windows
 4. **Aggregate:** Count, sum, average operations
 5. **Sink:** Kafka producer (analytics-results topic) + GCS
 - **Metrics Computed:**
 - Total orders per window
 - Revenue per window
 - Average order value
 - Top-selling products
 - Orders by region
 - Peak hour analysis
 - **Cluster Configuration:**
 - Master: n1-standard-2 (2 vCPU, 7.5GB RAM)
 - Workers: 2x n1-standard-2
 - Parallelism: 4
 - Checkpointing: Every 60 seconds to GCS
 - **Resource Allocation:**
 - Job Manager: 2GB heap
 - Task Manager: 4GB heap
 - Network buffers: 512MB
-

Service Interconnections & Communication Mechanisms

Detailed Communication Flow Diagram







Inter-Service Authentication Flow:



1. User Login:

Client □ API Gateway □ User Service
 POST /api/users/login {username, password}
 □ JWT Token (expires in 1 hour)

2. Authenticated Request:

Client □ API Gateway □ Order Service
 Headers: {Authorization: Bearer <JWT>}
 Order Service validates JWT
 □ Response with order data

3. Service-to-Service:

Order Service □ Product Service
 Headers: {Authorization: Bearer <SERVICE_TOKEN>}
 (Service tokens from Kubernetes secrets)

Communication Patterns in Detail

1. Synchronous Communication (HTTP REST) Pattern: Request-Response over HTTP/HTTPS

Use Cases: - Client □ API Gateway: All user-initiated actions - API Gateway □ Microservices: Request routing
 - Order Service □ Product Service: Stock validation - Order Service □ Payment Service: Payment processing -
 Order Service □ User Service: User validation

Service Discovery: - Uses Kubernetes DNS: `http://product-service.default.svc.cluster.local:80`
 - Short form: `http://product-service:80` - Environment variables for service URLs

Benefits: - Simple to understand and debug - Immediate response - Strong consistency

Challenges: - Tight coupling - Cascading failures - Increased latency

Mitigation: - Circuit breakers (future: Istio) - Timeouts and retries - Health checks - Fallback responses

2. Asynchronous Communication (Apache Kafka) Pattern: Event-driven messaging

Use Cases: - Order Service □ Notification Service: Order events - Order Service □ Analytics Service: Business intelligence - Payment Service □ Order Service: Payment status (future)

Kafka Configuration:

```
Cluster: MSK (3 brokers)
Topic: orders
  Partitions: 3
  Replication Factor: 2
  Retention: 7 days
  Compression: snappy
```

```
Topic: analytics-results
Partitions: 3
Replication Factor: 2
Retention: 30 days
Compression: gzip
```

Implementation: Order Service publishes events to Kafka with `acks='all'` for durability. Notification Service consumes with manual commits for guaranteed processing.

Benefits: - Loose coupling - Asynchronous processing - Event replay capability - Scalability - Fault tolerance

Challenges: - Eventual consistency - Message ordering - Duplicate handling

Mitigation: - Idempotent consumers - Message deduplication - Offset management - Dead letter queues

3. Database Access Patterns

Pattern: Direct database connections

User Service □ PostgreSQL: Uses SQLAlchemy ORM for database operations with session-based transactions

Product Service □ DynamoDB: Uses Boto3 SDK for NoSQL operations (`get_item`, `put_item`, `query` with GSI)

Connection Pooling: - PostgreSQL: pgbouncer (session pooling) - Max connections per service: 20 - Connection timeout: 30 seconds

4. Caching Strategy (Redis)

Pattern: Read-through cache with 1-hour TTL for product catalog, user sessions, and rate limiting

Service Resilience Patterns

Circuit Breaker (Planned) Uses pybreaker library with 5-failure threshold in 60-second window for service degradation

Retry with Exponential Backoff Implemented using tenacity library: 3 attempts with exponential wait (1s to 10s)

Timeout Configuration

- API Gateway □ Services: 30 seconds
- Service □ Service: 5 seconds
- Service □ Database: 10 seconds
- Kafka producer: 10 seconds
- Health checks: 3 seconds

Data Consistency Patterns

Strong Consistency (ACID Transactions) User, Order, and Payment services use PostgreSQL transactions with rollback on failure

Eventual Consistency (Event-Driven)

- Order events to notifications
- Order events to analytics
- Payment status updates

Strategy: - Idempotent operations - Event versioning - Compensating transactions for failures

Security in Service Communication

Mutual TLS (Future Enhancement) Service mesh (Istio) planned for strict mTLS enforcement between all services

API Authentication

- **External:** JWT tokens from User Service
- **Internal:** Service accounts with Kubernetes secrets

Network Policies Kubernetes NetworkPolicy configured to restrict traffic: Order Service only accepts ingress from API Gateway, egress to Product Service

Design Rationale & Decision Log

Architectural Decisions

Decision 1: Why Microservices Over Monolith? **Context:** Need to build scalable e-commerce platform

Options Considered: 1. Monolithic architecture 2. Microservices architecture 3. Modular monolith

Decision: Microservices Architecture

Rationale: - **Independent Scaling:** Order service experiences 10x traffic during sales, while user service remains steady. Microservices allow scaling only what's needed - **Technology Flexibility:** Analytics needs Java/Flink for stream processing, while REST APIs work well in Python. Microservices enable polyglot architecture - **Fault Isolation:** Payment service failure shouldn't crash product browsing. Microservices provide bulkheads - **Team Autonomy:** Different teams can own different services with independent release cycles - **Deployment Flexibility:** Can deploy order service updates without touching user service

Trade-offs Accepted: - Increased operational complexity (mitigated by Kubernetes) - Network latency between services (acceptable for our SLAs) - Eventual consistency challenges (managed with event sourcing) - Higher infrastructure costs (offset by independent scaling efficiency)

Decision 2: Why Apache Kafka for Event Streaming? **Context:** Need asynchronous communication between services

Options Considered: 1. Apache Kafka 2. RabbitMQ 3. AWS SQS 4. Direct HTTP callbacks

Decision: Apache Kafka (AWS MSK)

Rationale: - **Durability:** Messages persist for 7 days, enabling event replay and debugging - **Scalability:** Handles millions of events per second with partitioning - **Decoupling:** Order service doesn't know about consumers (notification, analytics) - **Stream Processing:** Native integration with Apache Flink for analytics - **Managed Service:** MSK eliminates operational overhead of running Kafka - **Ordering Guarantees:** Per-partition ordering crucial for order processing - **Multiple Consumers:** Both notification and analytics consume same events

Trade-offs Accepted: - Higher cost than SQS (\$300/month for MSK vs \$50 for SQS) - Learning curve for team - Eventual consistency model - More complex than simple queues

Why Not Alternatives: - RabbitMQ: Less suited for high-throughput streaming - SQS: No event replay, limited retention, not ideal for stream processing - HTTP callbacks: Tight coupling, retry complexity, no persistence

Decision 3: Why Multi-Cloud (AWS + GCP)? **Context:** Need to deploy application infrastructure

Options Considered: 1. Single cloud (AWS only) 2. Single cloud (GCP only) 3. Multi-cloud (AWS + GCP) 4. On-premises + cloud hybrid

Decision: Multi-Cloud (AWS Primary + GCP for Analytics)

Rationale: - **Best-of-Breed:** AWS excels at application hosting, GCP excels at data processing - **Risk Mitigation:** Not locked into single vendor; disaster recovery across clouds - **Cost Optimization:** GCP Dataproc more cost-effective for Flink than AWS EMR - **Learning Opportunity:** Team gains multi-cloud expertise - **Workload Isolation:** Analytics workload separated from transactional load

Trade-offs Accepted: - Increased operational complexity (two cloud consoles) - Cross-cloud networking costs - Skill requirement for two platforms - Inconsistent tooling and APIs

Implementation Strategy: - Keep compute and data gravity in AWS - Use GCP only for analytics workload - Kafka as integration point between clouds - Minimize cross-cloud data transfer

Decision 4: Why Kubernetes (EKS) Over ECS or Lambda? **Context:** Need container orchestration platform

Options Considered: 1. Amazon EKS (Kubernetes) 2. Amazon ECS (proprietary) 3. AWS Lambda (serverless) 4. EC2 instances directly

Decision: Amazon EKS (Kubernetes)

Rationale: - **Industry Standard:** Kubernetes is cloud-agnostic, transferable skills - **Rich Ecosystem:** Helm charts, operators, service meshes, monitoring tools - **Portability:** Can migrate to GKE or on-premises if needed - **Advanced Features:** HPA, network policies, custom schedulers, StatefulSets - **Community Support:** Massive community, extensive documentation - **GitOps Ready:** Native integration with ArgoCD for declarative deployments

Trade-offs Accepted: - Steeper learning curve than ECS - Control plane costs ($\$0.10/\text{hour} = \$73/\text{month}$) - More complex than serverless - Requires more operational expertise

Why Not Alternatives: - ECS: Vendor lock-in to AWS, less feature-rich, smaller community - Lambda: Cold starts, 15-minute timeout limit, not suitable for long-running processes - EC2: Manual scaling, no self-healing, higher operational burden

Decision 5: Why PostgreSQL (RDS) + DynamoDB? **Context:** Need database for different data patterns

Options Considered: 1. PostgreSQL for everything 2. DynamoDB for everything 3. MongoDB 4. PostgreSQL + DynamoDB (chosen)

Decision: PostgreSQL (RDS) for transactional + DynamoDB for catalog

Rationale:

PostgreSQL for Users/Orders/Payments: - **ACID Transactions:** Critical for financial data integrity - **Complex Queries:** JOIN operations for order items, user orders - **Relational Data:** Natural fit for user-order-payment relationships - **Referential Integrity:** Foreign keys ensure data consistency - **Mature Ecosystem:** ORMs, migration tools, backup solutions

DynamoDB for Products: - **High Throughput:** Product catalog read-heavy (90% reads) - **Flexible Schema:** Product attributes vary by category - **Predictable Performance:** Single-digit millisecond latency at scale - **Auto-Scaling:** Pay-per-request pricing scales with traffic - **No Maintenance:** Fully managed, no patches or upgrades

Trade-offs Accepted: - Two databases to manage instead of one - Different query patterns and SDKs - Cross-database joins not possible (mitigated by API composition)

Decision 6: Why GitOps with ArgoCD? **Context:** Need CI/CD deployment strategy

Options Considered: 1. kubectl apply in CI/CD pipeline 2. Helm direct installs 3. ArgoCD (GitOps) 4. Flux CD (GitOps alternative)

Decision: ArgoCD for GitOps

Rationale: - **Declarative:** Git is single source of truth for cluster state - **Auditability:** Every change tracked in Git history - **Easy Rollback:** git revert to previous working state - **Security:** No kubectl credentials in CI/CD; ArgoCD pulls changes - **Drift Detection:** Auto-sync corrects manual kubectl changes - **Multi-Cluster:** Can manage multiple Kubernetes clusters from one place - **Visualization:** UI shows deployment status and sync health

Trade-offs Accepted: - Additional component to run in cluster - Git becomes critical path (mitigated by multiple repos) - Learning curve for GitOps concepts

Why Not Alternatives: - kubectl in CI/CD: Requires credentials in pipeline, no drift detection - Helm direct: No automated sync, manual updates - Flux CD: ArgoCD has better UI and multi-tenancy support

Decision 7: Why Python Flask for Microservices? **Context:** Choose technology stack for REST APIs

Options Considered: 1. Python with Flask 2. Node.js with Express 3. Java with Spring Boot 4. Go with Gin

Decision: Python with Flask (except Analytics)

Rationale: - **Team Expertise:** Team already proficient in Python - **Fast Development:** Flask is lightweight and productive - **Rich Ecosystem:** SQLAlchemy, Kafka clients, AWS SDKs - **Readability:** Python code is maintainable and self-documenting - **Data Science Integration:** Easy to add ML features later - **Docker Support:** Official slim Python images are lightweight

Trade-offs Accepted: - Lower raw performance than Go or Java - GIL limitations for CPU-bound tasks (not an issue for I/O-bound APIs) - Async support less mature than Node.js

Why Analytics in Java?: - Apache Flink is Java-native - Better performance for stream processing - Mature Flink SDK in Java

Decision 8: Why Horizontal Pod Autoscaling (HPA)? **Context:** Need to scale applications based on load

Options Considered: 1. Static replica count 2. Horizontal Pod Autoscaler (HPA) 3. Vertical Pod Autoscaler (VPA) 4. Manual scaling

Decision: HPA for API Gateway and Order Service, Static for others

Rationale:

HPA for API Gateway & Order Service: - **Variable Load:** Traffic varies 10x between off-peak and sales events - **Cost Optimization:** Scale down during low traffic to save money - **Automatic:** No manual intervention during traffic spikes - **Reliable:** Kubernetes manages scaling decisions based on metrics

Static Replicas for User/Product/Payment/Notification: - **Predictable Load:** User CRUD operations are steady - **Simplicity:** No need for scaling complexity - **Faster Startup:** Always have capacity ready - **Cost:** Minimal cost difference (2-3 pods)

Configuration: - API Gateway: 2-10 replicas at 70% CPU - Order Service: 2-8 replicas at 70% CPU - Scale up fast (100% increase every 30s) - Scale down slow (50% decrease every 5min)

Trade-offs Accepted: - Metric collection overhead (Prometheus) - Potential overprovisioning during scale-up - Complexity in tuning thresholds

Decision 9: Why Prometheus + Grafana for Observability? **Context:** Need monitoring and alerting

Options Considered: 1. CloudWatch only 2. Prometheus + Grafana 3. Datadog (commercial) 4. New Relic (commercial)

Decision: Prometheus + Grafana + CloudWatch

Rationale: - **Open Source:** No per-host licensing costs - **Kubernetes Native:** Service discovery, pod metrics out-of-box - **Flexible Querying:** PromQL for custom queries and dashboards - **Rich Ecosystem:** Exporters for databases,

Kafka, etc. - **Alertmanager**: Flexible alerting with Slack/PagerDuty integration - **Grafana**: Beautiful dashboards, template variables, annotations

CloudWatch for: - AWS-specific metrics (EKS, RDS, MSK) - Log aggregation - Billing alerts

Trade-offs Accepted: - Need to run Prometheus in cluster - Data retention limits (30 days default) - Learning curve for PromQL

Why Not Commercial: - Datadog: \$15/host/month = \$500+/month for our scale - New Relic: Similar pricing - Budget-conscious for learning project

Infrastructure Design Rationale

Multi-AZ Deployment Decision: Deploy across 2 Availability Zones

Rationale: - **High Availability**: Survive single AZ failure - **Load Distribution**: Traffic balanced across zones -

Database Resilience: RDS Multi-AZ automatic failover

Cost Impact: - 2x data transfer costs between AZs - Acceptable for production-grade availability

VPC Design (10.0.0.0/16) **Decision**: Public subnets for EKS, private subnets for databases

Rationale: - **Security**: Databases not publicly accessible - **Compliance**: Industry best practice - **NAT Gateway**: Future for egress from private subnets

Subnet Allocation: - 10.0.0.0/24 - Public Subnet AZ1 (256 IPs) - 10.0.1.0/24 - Public Subnet AZ2 (256 IPs) - 10.0.10.0/24 - Private Subnet AZ1 (256 IPs) - 10.0.11.0/24 - Private Subnet AZ2 (256 IPs)

Resource Sizing EKS Node Group: - Instance Type: t3.medium (2 vCPU, 4GB RAM) - Min: 2 nodes, Max: 4 nodes - **Rationale**: - t3 burstable for variable workloads - medium size balances cost and capacity - Can run 8-12 pods per node (256Mi pods)

RDS Instance: - Instance Type: db.t3.micro (2 vCPU, 1GB RAM) - Storage: 20GB gp3 - **Rationale**: - Free tier eligible - Sufficient for development/demo - Can scale vertically if needed

Security Groups EKS Node Security Group: - Ingress: Port 443 from ALB - Egress: All (for pulling images, calling AWS APIs)

RDS Security Group: - Ingress: Port 5432 from EKS nodes only - No public internet access

MSK Security Group: - Ingress: Port 9092 from EKS nodes - Cross-cloud: Port 9092 from GCP IP range

Future Enhancements & Trade-offs

What We Would Add in Production

1. Service Mesh (Istio):

- Mutual TLS between services
- Advanced traffic management (canary, blue-green)
- Distributed tracing
- **Why not now**: Complexity for learning project

2. API Gateway Layer (Kong/Ambassador):

- Rate limiting per user
- API key management

- Request transformation
- **Why not now:** Kubernetes Ingress sufficient for demo

3. Caching Layer (Redis):

- Product catalog caching
- Session storage
- Rate limit counters
- **Why not now:** Database performance adequate currently

4. CDN (CloudFront):

- Static asset delivery
- Global edge caching
- DDoS protection
- **Why not now:** No static assets currently

5. Secrets Management (Vault/Secrets Manager):

- Dynamic secrets rotation
- Centralized secret storage
- Audit logging
- **Why not now:** Kubernetes Secrets sufficient for demo

6. Service Mesh Observability:

- Distributed tracing (Jaeger)
- Service topology visualization
- Request flow analysis
- **Why not now:** Prometheus metrics sufficient

7. Multi-Region Deployment:

- Active-active across regions
 - Global load balancing
 - Data replication
 - **Why not now:** Regional deployment sufficient, cost prohibitive
-

Cost-Benefit Analysis

Current Monthly Costs (estimated): - EKS Control Plane: \$73 - EC2 Instances (2x t3.medium): \$60 - RDS (db.t3.micro): \$15 - MSK (kafka.t3.small): \$300 - DynamoDB: \$10 (on-demand) - Data Transfer: \$20 - **Total:** ~\$478/month

Cost Optimization Strategies: 1. Reserved Instances: Save 30-40% on EC2/RDS 2. Spot Instances: Use for dev/test environments 3. Auto-scaling: Scale down during off-hours 4. Data lifecycle: Move old data to S3 Glacier

Cost vs. Alternative Architectures: - Serverless (Lambda + API Gateway): ~\$200/month but cold starts - Monolith on single EC2: ~\$50/month but no scaling - Managed containers (ECS Fargate): ~\$350/month but less features

Value Delivered: - Production-grade architecture - Skills development in modern DevOps - Portfolio showcase - Real-world experience with cloud-native patterns

Infrastructure as Code

Terraform Modules

1. **VPC Module:** Network isolation, public/private subnets
2. **EKS Module:** Kubernetes cluster with node groups
3. **RDS Module:** MySQL database with security groups
4. **MSK Module:** Kafka cluster configuration
5. **Lambda Module:** S3-triggered function
6. **GCP Module:** Dataproc cluster and Cloud Storage

Benefits

- **Reproducibility:** Spin up identical environments
 - **Version control:** Infrastructure changes tracked
 - **Collaboration:** Team can review infrastructure changes
 - **Automation:** CI/CD can apply infrastructure changes
-

Scalability & Resilience

Horizontal Pod Autoscaling (HPA)

- **API Gateway:** 2-10 pods based on CPU (70%)
- **Order Service:** 2-8 pods based on CPU (70%)
- **Triggers:** CPU/memory utilization

Database Scaling

- **RDS:** Vertical scaling (instance size)
- **DynamoDB:** Auto-scaling (pay-per-request)

High Availability

- **Multi-AZ:** RDS and EKS nodes across 2 AZs
 - **Load Balancing:** AWS ELB for API Gateway
 - **Health Checks:** Kubernetes liveness/readiness probes
-

Observability

Metrics (Prometheus + Grafana)

- Request rate, latency, error rate per service
- Kubernetes cluster health
- Database connections
- Kafka lag

Logging (EFK Stack)

- Centralized log aggregation
- Application logs from all pods
- Searchable via Kibana

Tracing (Future)

- Distributed tracing with OpenTelemetry
 - End-to-end request tracking
-

Security Considerations

1. **Network Isolation:** Private subnets for databases
2. **IAM Roles:** Service accounts with least privilege
3. **Secrets Management:** Kubernetes secrets for credentials
4. **TLS:** HTTPS for public endpoints
5. **Security Groups:** Firewall rules limiting access

Cost Optimization

Free Tier Usage

- **RDS**: db.t3.micro (free tier eligible)
- **EKS**: Control plane free first 30 days
- **Lambda**: 1M requests free
- **S3/DynamoDB**: Generous free tiers

Cost Management

- **Auto-scaling**: Pay only for what you use
- **Spot instances**: (Future) For non-critical workloads
- **Monitoring**: CloudWatch to track spending

Estimated Monthly Cost: \$50-100 (mostly MSK Kafka)

Future Enhancements

1. **Service Mesh** (Istio): Traffic management, security
 2. **API Gateway** (Kong): Rate limiting, authentication
 3. **Caching** (Redis): Reduce database load
 4. **CDN** (CloudFront): Static asset delivery
 5. **Multi-region**: Global deployment
-

Conclusion

This architecture demonstrates: - Modern cloud-native patterns - Multi-cloud deployment - Event-driven design - GitOps workflows - Comprehensive observability - Production-ready practices

All while remaining cost-effective and manageable.